Lecture 01

Bellman equation : $Q(s,a) = r_t + \gamma \max_{a'} Q(s',a')$

actions; $0 \le \gamma \le 1$; current state; next step reward

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

$$= r_t + \gamma G_{t+1}$$

	a_0	a_1	...	a_m
s_0	$Q(s_0,a_0)$			
s_1				
s_2				
⋮				
s_n				

$Q(s_k, a_l)$

① Scale is large s, a

② Convergence issues because of lack of time

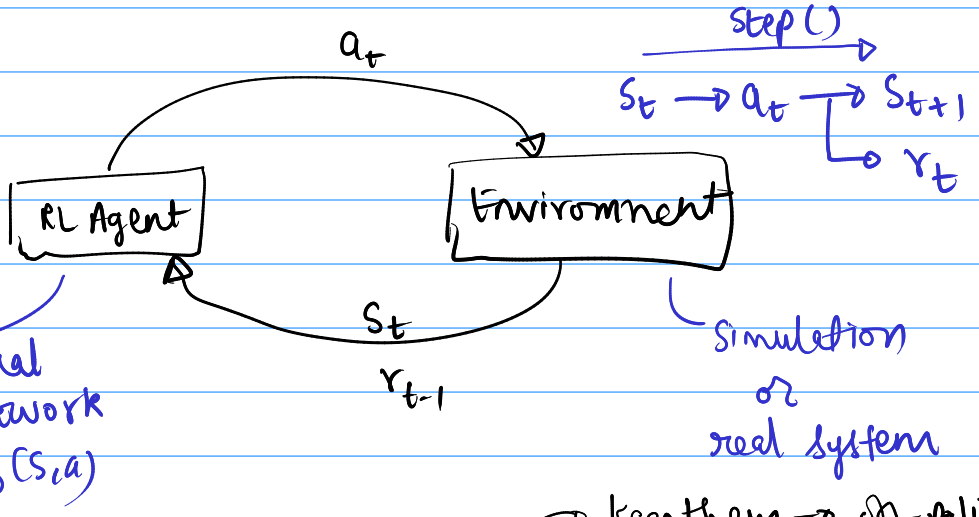

"sample" : $[s_t, a_t, r_t, s_{t+1}]$ $\forall t$

↳ keep them → off-policy

↳ throw away → on-policy

$$Q(s_t, a_t) = r_t + \gamma \max_{a_{t+1}} Q(s_{t+1}, a_{t+1})$$

Predictions or inference; target

Error = MSE $= [Q(s_t, a_t) - r_t - \gamma \max Q(s_{t+1}, a_{t+1})]^2$

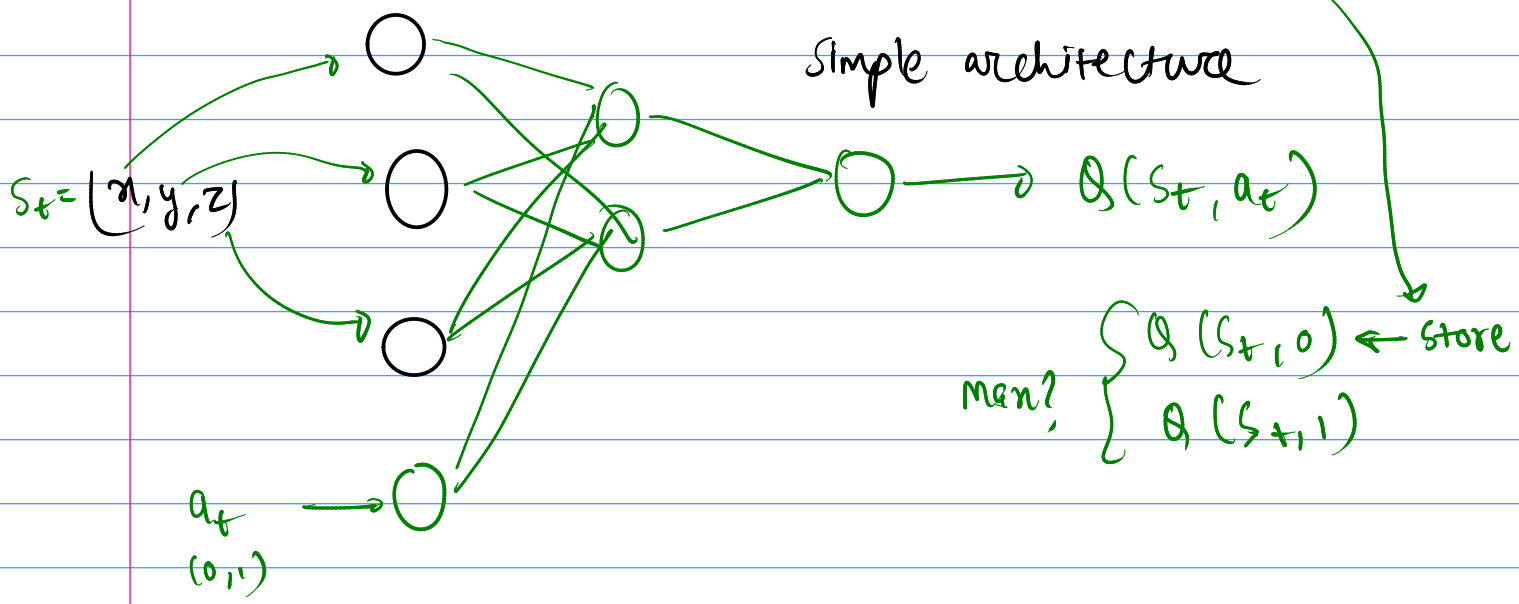

Practical architecture

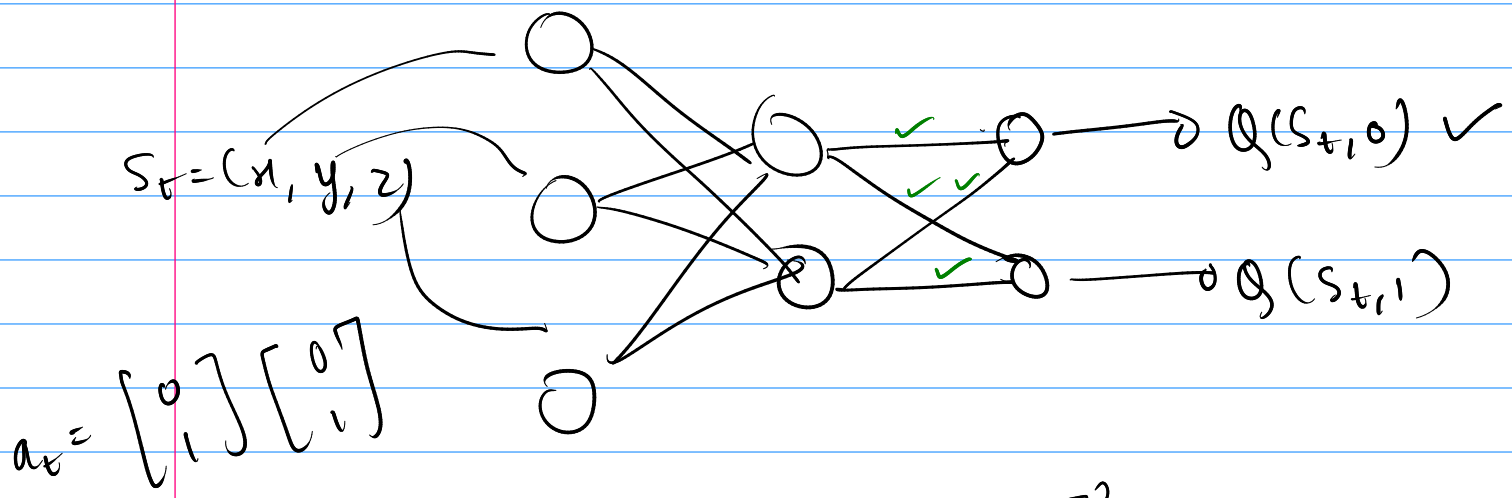

$$[Q(s_t, 0) - r + \gamma Q(s_{t+\ldots})]^2 + 0$$

"Vanilla DQN"

Issues

① Bootstrapping
② Stabilisation
③ $\Delta Q(s,a)$ implemented only by last layer
④ Off/on policy decisions

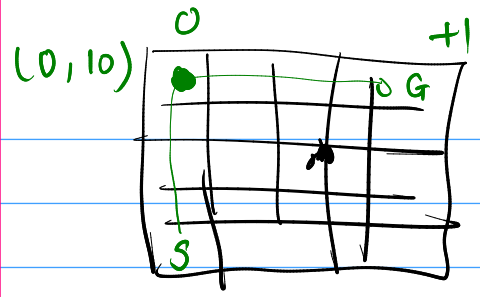

(x,y) → 0 → $Q(S_{t+1}, a_{t+1})$

(S_t, a_t, r_t, S_{t+1})

$[\ S_t, (10,10), a_t, r_t, S_{t+1}$
$\vdots$
$]$ 0

$[\ \boxed{S_t, (0,10),} a_t, r_t, S_{t+1}$
$\vdots$
$(0,10)\ \ (0,10), a_t, 1, S_{t+1}\ ]$

within agent

① choose_action ⟶ state as input
compute $Q(S,a)$ by inference on NN
G-greedy action

② save to memory ⟶ S_t, a_t, r_t, S_{t+1}

③ train ⟶ sample from memory
compute loss
update NN

$$Q(S,a) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots$$
$$= \underbrace{r_t + \gamma r_{t+1}} + \gamma^2 [\underbrace{r_{t+2} \cdots}_{G_{t+2}}$$