

Coac / never-give-up Public[Code](#) [Issues 1](#) [Pull requests](#) [Actions](#) [Projects](#) [Security](#) [...](#)

main ▾

...

[never-give-up](#) / [train.py](#) / [Jump to ▾](#)

Coac Add working config

[History](#)

1 contributor

161 lines (124 sloc) | 5.11 KB

...

```
1 import numpy as np
2 import torch
3 import torch.optim as optim
4 import wandb
5 from gym import Wrapper
6 from gym_maze.envs.maze_env import MazeEnvSample5x5
7
8 from config import config
9 from embedding_model import EmbeddingModel, compute_intrinsic_reward
10 from memory import Memory, LocalBuffer
11 from model import R2D2
12
13 def get_action(state, target_net, epsilon, env, hidden):
14     action, hidden = target_net.get_action(state, hidden)
15
16     if np.random.rand() <= epsilon:
17         return env.action_space.sample(), hidden
18     else:
19         return action, hidden
20
21 def update_target_model(online_net, target_net):
22     target_net.load_state_dict(online_net.state_dict())
23
24 class Maze(Wrapper):
25     def step(self, action: int):
26         obs, rew, done, info = super().step(["N", "E", "S", "W"][action])
27         self.set.add((obs[0], obs[1]))
28         if rew > 0:
29             rew = 10
```

 π_t in paperfrom R2D2 code, produced by
LSTM
handles
POMDP→ ϵ -greedy← straight copy of
online network
weights

move the agent

```

33         return obs / 10, rew, done, info
34         state reward
35     def reset(self):
36         self.set = set()
37         return super().reset()
38
39
40 def main():
41     env = Maze(MazeEnvSample5x5())
42
43     torch.manual_seed(config.random_seed)
44     env.seed(config.random_seed)
45     np.random.seed(config.random_seed)
46     env.action_space.seed(config.random_seed)
47
48     wandb.init(project="ngu-maze", config=config.__dict__)
49
50     num_inputs = env.observation_space.shape[0]
51     num_actions = env.action_space.n
52     print("state size:", num_inputs)
53     print("action size:", num_actions)
54     see model.py
55     online_net = R2D2(num_inputs, num_actions)
56     target_net = R2D2(num_inputs, num_actions)
57     update_target_model(online_net, target_net)
58     embedding_model = EmbeddingModel(obs_size=num_inputs, num_outputs=num_actions)
59     embedding_loss = 0
60
61     optimizer = optim.Adam(online_net.parameters(), lr=config.lr)
62
63     online_net.to(config.device)
64     target_net.to(config.device)
65     online_net.train()
66     target_net.train()
67     memory = Memory(config.replay_memory_capacity)
68     epsilon = 1.0
69     steps = 0
70     loss = 0
71     local_buffer = LocalBuffer()
72     sum_reward = 0
73     sum_augmented_reward = 0
74     sum_obs_set = 0
75
76     for episode in range(30000):
77         done = False
78         state = env.reset()
79         state = torch.Tensor(state).to(config.device)
80
81         hidden = (
82             torch.Tensor().new_zeros(1, 1, config.hidden_size),
83             torch.Tensor().new_zeros(1, 1, config.hidden_size),

```

termination flag

state reward

Note for reproducibility
Set random seeds within
every class

Same weights at the start

embedding-model.py → see left side of fig 1
in paper

number of episodes

```

84         )
85
86         episodic_memory = [embedding_model.embedding(state)]
87
88         episode_steps = 0
89         horizon = 100
90         while not done:
91             steps += 1
92             episode_steps += 1
93
94             action, new_hidden = get_action(state, target_net, epsilon, env, hidden)
95
96             next_state, env_reward, done, _ = env.step(action)
97             next_state = torch.Tensor(next_state)
98
99             augmented_reward = env_reward
100             if config.enable_ngu:
101                 next_state_emb = embedding_model.embedding(next_state)
102                 intrinsic_reward = compute_intrinsic_reward(episodic_memory, next_state_emb)
103                 episodic_memory.append(next_state_emb)
104                 beta = 0.0001
105                 augmented_reward = env_reward + beta * intrinsic_reward
106
107             mask = 0 if done else 1
108
109             local_buffer.push(state, next_state, action, augmented_reward, mask, hidden)
110             hidden = new_hidden
111             if len(local_buffer.memory) == config.local_mini_batch:
112                 batch, lengths = local_buffer.sample()
113                 td_error = R2D2.get_td_error(online_net, target_net, batch, lengths)
114                 memory.push(td_error, batch, lengths)
115
116             sum_reward += env_reward
117             state = next_state
118             sum_augmented_reward += augmented_reward
119
120             if steps > config.initial_exploration and len(memory) > config.batch_size:
121                 epsilon -= config.epsilon_decay
122                 epsilon = max(epsilon, 0.4)
123
124                 batch, indexes, lengths = memory.sample(config.batch_size)
125                 loss, td_error = R2D2.train_model(online_net, target_net, optimizer,
126                                                  # RL training
127                                                  if config.enable_ngu:
128                                                      embedding_loss = embedding_model.train_model(batch)
129                                                      # training embedding
130              memory.update_priority(indexes, td_error, lengths)
131
132             if steps % config.update_target == 0:
133                 update_target_model(online_net, target_net)
134

```

replace

running in NGU mode

see embedding-model.py

RL training

training embedding

slow update after fixed number of updates of online model

```
135         if episode_steps >= horizon or done:
136             sum_obs_set += len(env.set)
137             break
138
139         if episode > 0 and episode % config.log_interval == 0:
140             mean_reward = sum_reward / config.log_interval
141             mean_augmented_reward = sum_augmented_reward / config.log_interval
142             metrics = {
143                 "episode": episode,
144                 "mean_reward": mean_reward,
145                 "epsilon": epsilon,
146                 "embedding_loss": embedding_loss,
147                 "loss": loss,
148                 "mean_augmented_reward": mean_augmented_reward,
149                 "steps": steps,
150                 "sum_obs_set": sum_obs_set / config.log_interval,
151             }
152             print(metrics)
153             wandb.log(metrics)
154
155             sum_reward = 0
156             sum_augmented_reward = 0
157             sum_obs_set = 0
158
159
160 if __name__ == "__main__":
161     main()
```

*useful
for
tracking
if you
haven't
tried it*